

Amazon Sagemaker is an umbrella of services that AWS provides for Machine Learning.

A brief overview of main products include:

Sagemaker Studio

Sagemaker Studio is a cloud-based IDE, centered around Jupyter Lab but focused specifically on Machine Learning practitioners in mind. It allows you to complete the whole lifecycle of machine learning, from creating datasets, performing EDA, to training and hosting your models. Because it's cloud-first, you can share and collaborate easily among peers.

Ground Truth

Ground Truth is a service for annotating data. Today, many ML models require massive amounts of labeled data, and Ground Truth provides the infrastructure and processes to complete that end to end.

Data Wrangler

Data Wrangler allows you to perform exploratory data analysis, feature engineering your data and visualize insights all within the Sagemaker Studio interface.

Feature Store

Feature Store manages all of your features, providing a centralized location to store, update, share, and obtain features.

Training/Pipelining/Endpoints

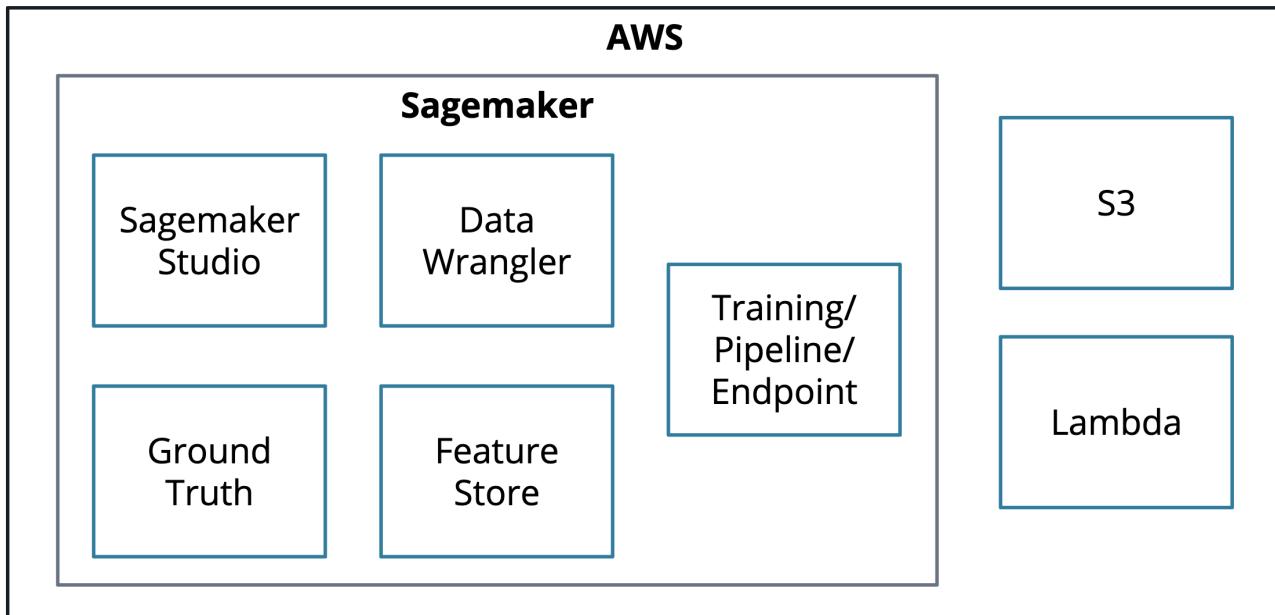
There are also several processes that combined include the training, pipelining, and serving of models to endpoints. These are more or less EC2 instances, general compute from AWS, but managed solely by Sagemaker.

S3

S3 is an object store. It allows you to store almost any object needed for machine learning. That includes datasets, model artifacts, logs, and more.

Lambda

AWS Lambda is serverless computing, which means they are tiny functions or single-purpose code which can help glue together pieces of the machine learning lifecycle.



Sagemaker Services Overview

Additional Resources

- If you want to learn more about Sagemaker, visit the [features](#) page on AWS.